



Semi-automatic literature reviews

Sam Dornan

What exactly is a semi-automatic literature review?

Using machine learning to infer the relationship between words in a corpus

- Can review larger volumes of text
- Find relationships between disparate disciplines and fields of study
- Already used across material sciences

Unsupervised word embeddings capture latent knowledge from materials science literature

[Vahe Tshitoyan](#) ✉, [John Dagdelen](#), [Leigh Weston](#), [Alexander Dunn](#), [Ziqin Rong](#), [Olga Kononova](#), [Kristin A. Persson](#), [Gerbrand Ceder](#) ✉ & [Anubhav Jain](#) ✉

[Nature](#) **571**, 95–98 (2019) | [Cite this article](#)

69k Accesses | **381** Citations | **1899** Altmetric | [Metrics](#)

How would it be done?

Social sciences is a part of the world as much as the world is a part of the social sciences



Large language models such as GPT and T-5 can “understand” the world.



Open-source programmes allow models to be tailored to a specific task.



Tailor a language model to the social sciences

The dataset

Easy to access

Easy to work with

Small-fits in memory

Tailored

Wisdom of the crowds

The screenshot displays the Web of Science interface. At the top, the Clarivate logo is on the left, and 'English' and 'Products' are on the right. Below the logo, the 'Web of Science' title and a search bar are visible. The search results show '170 results from Web of Science Core Collection for: Sociology (Topic)'. A sidebar on the left contains a 'MENU' icon and a list of filters: 'Highly Cited Papers' (170), 'Hot Papers' (4), 'Review Article' (30), 'Open Access' (92), and 'Associated Data' (8). The main results area shows a list of publications, with the first one being 'For a Sociology of Expertise: The Social Origins of the Autism Epidemic' by Eyal, G. (2013). The article is from the 'AMERICAN JOURNAL OF SOCIOLOGY' and has 247 citations and 167 references. The interface also includes an 'Export' dropdown menu with options like 'EndNote online', 'Excel', and 'Printable HTML file'.

The manual part?

Garbage-in; garbage-out

Preprocessing/Post-processing

Where's the data from?

What's in the data?

What language model do we use?

Why?

The screenshot displays the Web of Science interface. At the top, the Clarivate logo is on the left, and 'English' and 'Products' are on the right. The main header shows 'Web of Science' and a search bar. Below the search bar, the results are for 'Sociology (Topic)' with 170 results from the Web of Science Core Collection. A sidebar on the left contains a menu with icons for home, search, and other functions. The main content area shows the search results, including a list of publications. The first result is 'For a Sociology of Expertise: The Social Origins of the Autism Epidemic' by Eyal, G., published in the American Journal of Sociology in 2013. The article has 247 citations and 167 references. A dropdown menu is open, showing export options: EndNote online, EndNote desktop, Plain text file, RIS (other reference software), BibTeX, Excel, Tab delimited file, Printable HTML file, InCites, Email, Fast 5000, and More Export Options. The interface also includes a 'Refine results' section with a search bar and a 'Filter by Marked List' section with various filters like 'Highly Cited Papers', 'Hot Papers', 'Review Article', 'Open Access', and 'Associated Data'.

The Model

Pretrained

Small – fits in memory

Only compatible with T-5 by Raffel et al (2020).

LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale

Tim Dettmers^{λ*} Mike Lewis[†] Younes Belkada^{§‡} Luke Zettlemoyer^{†λ}

University of Washington^λ
Facebook AI Research[†]
Hugging Face[§]
ENS Paris-Saclay[‡]


	GLUE	CNN3M	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Baseline average	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Baseline standard deviation	0.235	0.065	0.343	0.416	0.112	0.090	0.108
No pre-training	66.22	17.60	50.31	53.04	25.86	39.77	24.04

Table 1: Average and standard deviation of scores achieved by our baseline model and training procedure. For comparison, we also report performance when training on each task from scratch (i.e. without any pre-training) for the same number of steps used to fine-tune the baseline model. All scores in this table (and every table in our paper except Table 14) are reported on the validation sets of each data set.

The inputs and outputs - Further research



Semantic Search
Programme



Semantic Network Graph

References

DETTMERS, T., LEWIS, M., BELKADA, Y. & ZETTLEMOYER, L. 2022. LLM.int8 (): 8-bit Matrix Multiplication for Transformers at Scale. *arXiv preprint arXiv:2208.07339*.

RAFFEL, C., SHAZEER, N., ROBERTS, A., LEE, K., NARANG, S., MATENA, M., ZHOU, Y., LI, W. & LIU, P. J. 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *J. Mach. Learn. Res.*, 21, 1-67.

TSHITOYAN, V., DAGDELEN, J., WESTON, L., DUNN, A., RONG, Z., KONONOVA, O., PERSSON, K. A., CEDER, G. & JAIN, A. 2019. Unsupervised word embeddings capture latent knowledge from materials science literature. *Nature*, 571, 95-98.

YOSHITAKE, M., SATO, F., KAWANO, H. & TERAOKA, H. 2022. MaterialBERT for natural language processing of materials science texts. *Science and Technology of Advanced Materials: Methods*, 2, 372-380.